DWAN Command Center: the eight tabs explained in plain language

Oct 10, 2026 · @Ram

Start here: eight questions, one Monday morning

The eight tabs are the eight questions a busy executive asks about cyber risk, in the order they naturally come up. The prototype shows one sample company, Al Dhabi Holdings, on 5 September 2026. Every figure is made-up sample data.

Think of a health check-up. First the doctor tells you if you are fine overall (tab 1). Then which organ is weakest (2), whether the tests themselves can be trusted (3), what needs treatment today (4), what a treatment would change (5), whether last month's treatment worked (6), how you compare with the last visit (7), and finally the report you hand to the family (8).

To make each tab concrete, we follow **Sara**, the CXO of Al Dhabi Holdings, through one Monday. Sara is a made-up character; the numbers she sees are the ones on the real screens.

#	Tab	The question	The sample answer
1	Board headline	Are we OK?	No. Risk is High (50.8), because 2 domains are Critical.
2	Domain posture	Where are we weakest?	5 of 12 domains are outside what the board accepted.
3	Trust panel	Can I believe the numbers?	35 of 39 data feeds are fresh. Confidence is medium (72%).
4	Decision queue	What needs my decision today?	4 decisions, ranked. The top one is the payment gateway.
5	What-if	What should I fund?	Fixing ransomware alone does not change the headline.
6	Closure and acceptance	Did it work?	2 fixes are reported but not yet proven. 1 accepted risk is about to expire.
7	What changed	What moved since the last board meeting?	Headline unchanged. 4 domains worse, 5 better.

8 Board pack	What do I send the board?	A draft pack, where every sentence names its data source.

The bar at the top lets you view the same data as a Board member, CXO, CISO, Risk lead or Analyst. The numbers do not change; the buttons do. A board member can look and comment, but cannot approve anything.

1. Board headline: "Are we OK?"

This tab gives one honest verdict, and for Al Dhabi it is High, not Medium, even though the average looks gentler.

In plain words. Imagine a building with 12 rooms. Eleven are fine and one is on fire. Nobody would call the building "mostly fine". The headline works the same way: if any domain is Critical, the overall verdict cannot be better than High. This is called the floor rule, and the screen says when it was applied.

What you see on the screen

- A big badge, **High**, with the score **50.8** out of 100. Higher means more risk.
- The spread of the 12 domains: 2 Critical, 2 High, 4 Medium, 4 Low.
- **Confidence 72% (Medium)**: how much of this number rests on fresh, well-covered data.
- **5 of 12 domains are outside the board's appetite**, meaning above the limit the board said it would tolerate.
- **1 of 2 accepted risks expires within 30 days.**
- A money box: **AED 2.9M expected loss per year**, and **AED 10.07M in a bad year** (1 in 10).
- "Where to look first": Attack Surface and VAPT (76), SOC / MDR (72), Third-Party Risk (65) and Ransomware Readiness (55.2).

Real-time example. It is 8:30 on Monday and Sara has two minutes before a call with the CEO. She opens the first tab and reads: "Two domains are Critical and need decisions this week." She does not need to open a single scanner console. On the call she says: "We are at High. Modelled loss is about AED 2.9M in a normal year and about AED 10M in a bad one. Two things need your sign-off this week."

The money figure is like flood insurance. In most years nothing big happens, so the expected cost is small. The "bad year" number tells you what a rare, large event would cost.

Watch out for. The screen quotes two Medium figures: 42.8 as the plain average of the 12 raw scores, and 50.8 as the weighted, coverage-adjusted score. Both are Medium, but 50.8 is the one that drives the headline, so it is the one to quote.

2. Domain posture: "Where are we weakest?"

This tab lists all 12 risk domains side by side, and shows that 5 of them are outside the limit the board agreed to accept.

In plain words. A domain is one area of security, such as email, cloud or ransomware recovery. Each gets a score from 0 to 100, where higher means more risk. Think of 12 fuel gauges on a dashboard, each with a red line marking how low the board is willing to let it go. Here the line is the board's **appetite**: the maximum score it will tolerate for that domain.

How to read each row

- **Adjusted score**: the score after allowing for what the tools cannot see.
- **Raw**: the score from the data the tools did see.
- **Coverage**: the share of assets the tools actually watch.
- **Confidence**: how much to trust the row. A hatched bar means under 80%.
- **Appetite**: a white tick on the bar, plus a label: Within, Approaching or Outside.
- **6-month trend**: a small line, with arrows and signs as well as colour.
- **CSF 2.0**: which functions of the NIST Cybersecurity Framework 2.0 the domain supports, for example Detect and Respond.

Real-time example: the plant that looks safe. OT / ICS Security (the factory-floor systems) has a raw score of only 17, which looks excellent. But coverage is 40%: the tools watch just 4 in 10 plant assets. DWAN assumes the unwatched 6 are at least a 60. So the adjusted score is 0.4 × 17 + 0.6 × 60 = **42.8, Medium**. Without this step the plant would look as safe as a locked room that nobody has checked.

Real-time example: the one that is truly bad. PCI DSS / Data Security (card and personal data) scores 51 against a limit of 35, and coverage is 90% with high confidence. Sara can take this one at face value: it is a real problem, not a data gap. She taps the row and a side panel shows the evidence behind it.

Real-time example: the one to watch. AI Security and Governance scores 53.5 against a limit of 55, so it is "Approaching". Confidence is only 25%, because 2 of its 4 feeds are stale. Sara asks her team to fix the feeds before anyone trusts that number.

Watch out for. A low score with low coverage is not good news. Look for hatched bars first.

3. Trust panel: "Can I believe the numbers?"

This tab shows how much of every score rests on fresh, complete data, and for Al Dhabi the answer is 35 of 39 feeds fresh and a medium confidence of 72%.

In plain words. A bathroom scale is only useful if it is switched on, calibrated and you actually stood on it. Every score in DWAN comes from tools that send data in "feeds" (a feed is one tool's regular delivery of data). If a feed goes quiet, or the tool only watches part of your estate, the score is less reliable. This tab says so, instead of hiding it.

The formula, in words. Confidence = fresh feeds ÷ total feeds × coverage. Stale feeds and blind spots both pull it down. They are never counted as good news.

What you see on the screen

- **Feeds fresh: 35/39.** Four feeds are overdue.
- **Weighted coverage: 78%.** On average, tools see about 8 of every 10 assets.
- **Double-counting check: Pass.** No tool carries weight in two domains, so no score is inflated by counting one product twice.
- **Policy v1, draft, awaiting CISO sign-off.** The weights behind the scores are not yet approved, so the numbers are provisional.
- A table with fresh/total feeds, coverage, confidence and a plain-English reading for each domain, for example "2 of 4 feeds stale".

Real-time example: the quiet backup tool. On Friday evening the connector for the backup platform stops sending data after a software update. Over the weekend nothing looks wrong, because the old figures are still on screen. On Monday the Trust panel shows Ransomware Readiness as "1 of 4 feeds stale", and its confidence drops to 60% (3 of 4 feeds × 80% coverage). Sara's analyst sees it in minutes and restarts the connector.

Real-time example: the weakest row. AI Security and Governance shows 2 of 4 feeds stale and 50% coverage, so its confidence is 2/4 × 0.5 = **25%**. That is why its 53.5 on the previous tab is drawn hatched. Sara would not take a board decision on that number yet.

Watch out for. High confidence does not mean low risk. It only means the number is well supported. A domain can be Critical with 100% confidence, and that is the worst case to ignore.

4. Decision queue: "What needs my decision today?"

This tab turns thousands of security findings into four decisions, ranked, each with an owner, a deadline and a recommended action.

In plain words. A hospital triage nurse does not treat patients in the order they arrive. She ranks them by how dangerous the problem is, how fragile the patient is, and how exposed they are. DWAN does the same with findings. The **priority** is exploitability × how important the asset is × how exposed it is. It is shown out of 10.

Two terms you will see on the cards:

- **On CISA KEV**: the US government's list of weaknesses already being used by real attackers. Being on it is the strongest warning sign there is.
- **EPSS 0.84**: a 0 to 1 estimate of how likely a weakness is to be attacked soon. 0.84 means very likely.

What each card shows. The domain, the problem in one line, priority, the owner, the deadline (SLA), the asset affected, the modelled loss (expected and bad-year), a recommended action, and four buttons: **Approve**, **Accept risk**, **Defer** and **Evidence**.

Real-time example: the payment gateway. Card 1 reads "Two critical, internet-facing CVEs on the customer-facing payment gateway". It scores **10 out of 10**: it is on the KEV list, EPSS is 0.84, the asset is a crown jewel (the system the business cannot lose), and it faces the internet. The owner is the Head of Payments Engineering, the deadline is 7 days, and the modelled loss is AED 0.61M a year, rising to AED 2.75M in a bad year. The recommendation: approve an emergency patch window outside the normal change cycle. Sara presses **Approve**. The decision is recorded with her name and the time, and the owner's 7-day clock starts.

Real-time example: a quieter one. Card 2 says only 7 of 28 recovery tests met their target. That is a 25% pass rate, like fire drills where three in four fail to empty the building in time. It ranks 7 out of 10, with a 90-day deadline, and the recommendation is to fund a recovery-testing programme. Sara does not approve it yet. She opens the What-if tab first (tab 5).

Real-time example: when to say "accept". If a finding cannot be fixed now, Sara can choose **Accept risk**. The system then requires a reason, an approver and an expiry date, so it cannot be quietly forgotten (see tab 6).

Watch out for. A board member sees the same queue but cannot press the buttons. Only the CXO, CISO and risk roles can act.

5. What-if: "What should I fund?"

This tab lets you test a fix on paper before spending money, and it shows that fixing ransomware alone will not change the board headline.

In plain words. It is a flight simulator for budgets. You move a slider, such as "improve backups", and the page recalculates the score, the band and the expected loss as if the fix were done. Nothing here changes the live score, and nothing is saved unless you press **Save scenario**.

What you see on the screen

- Four sliders, one per ransomware control, each showing the remaining **gap** (how far from fully protected): EDR / XDR 52%, SIEM / SOAR 30%, Immutable backup 50%, DR / recovery testing 75%. Each also shows its weight, meaning how much it counts.
- Four presets: **Today**, **A. Immutable backup**, **B. Recovery testing**, **C. Both**.
- An optional cost box in AED millions.
- A before-and-after table: ransomware raw and adjusted score, band, enterprise score and headline, and expected and bad-year loss.

Real-time example: the CFO's question. The CFO asks, "If we pay for both, what do we get?" Sara presses **C. Both**. These are the sample results.

Scenario	Ransomware score (adjusted)	Band	Expected loss a year
Today	55.2	High	AED 1.78M
A. Immutable backup	47.2	Medium	AED 1.52M
B. Recovery testing	45.6	Medium	AED 1.47M
C. Both	37.6	Medium	AED 1.21M

Choosing C moves ransomware from High to Medium and cuts expected loss by about AED 0.57M a year. If the programme costs AED 1.5M (an illustrative figure Sara types into the cost box), the saving alone takes about 2.6 years to repay it. That is a conversation a budget committee can have.

The surprise. Look at the enterprise headline row: it stays **High**. Two other domains are still Critical, so the floor rule holds the headline up. The screen says so plainly. Sara learns that this investment is worth it for ransomware, but it will not by itself change what the board reads. The Critical payment gateway and SOC findings matter more for the headline.

Watch out for. The loss estimate here is a simple scaling: expected loss moves in proportion to the adjusted score. It is a good way to compare options, not a calibrated forecast. Replace it with the client's own loss model before quoting the money in a board paper.

6. Closure and acceptance: "Did it work?"

This tab separates "someone says it is fixed" from "a scan proves it is fixed", and it keeps every knowingly accepted risk on a short leash.

In plain words. When a mechanic says your car is repaired, you still take it for a test drive. A finding in DWAN is the same: the team reports a fix, and only a fresh scan can close it. And when the business chooses to live with a risk, it is like a parking permit: it names who approved it, why, and the date it runs out.

The two halves of the screen

1. **Verified closure.** Fixes waiting for proof, each with a **Run verification scan** button.
2. **Risk-acceptance register.** Risks the business knowingly accepted, each with a reason, an approver and an expiry, plus **Renew 90 days** and **Revoke and reopen** buttons.

Real-time example: two fixes, two outcomes. Two fixes are reported but not verified: a High vulnerability on an internal file server, and an exposed admin page on a test host. Sara's risk lead presses **Run verification scan** on each. In the sample, the first scan comes back clean, and the finding becomes **Verified closed**. The second scan still finds the admin page open, so the finding stays open and the delay keeps counting against its owner's deadline. Without this tab, both would have been marked "done" on the word of the person who did the work. The third item, outdated TLS on the partner portal, is already verified closed.

Real-time example: the permit about to expire. The register holds two accepted risks.

Accepted risk	Why it was accepted	Approver	Expires
Legacy HR portal without MFA	Retirement planned, and a network restriction compensates for now	CFO	29 Sep 2026 (24 days)
Unsupported OT historian software	Vendor replacement scheduled next quarter	COO	15 Nov 2026 (71 days)

The HR portal permit is nearly up, so it also shows on the Board headline and in the board pack. If the replacement is on track, Sara presses **Renew 90 days**. If not, **Revoke and reopen** puts it back in the decision queue.

Rules built in. The approver must be a senior executive, and never the person who owns the finding. The reason cannot be blank. An expiry is mandatory, with a shorter limit for Critical findings.

Watch out for. Accepted risks still count in the score. Accepting a risk is a decision to live with it, not a way to make it disappear from the numbers.

7. What changed: "What moved since the last board meeting?"

This tab compares today with the June board snapshot, and for Al Dhabi the headline is unchanged at High, with 4 domains worse and 5 better.

In plain words. It is a before-and-after photo taken with the same camera. The scoring method and the policy are identical in both columns, so any difference is real change in the company's security, not a change in how it is measured. That matters because board members often ask, "Did the score move because we improved, or because you changed the maths?"

What you see on the screen

- **Biggest movers**, each as old score → new score with an arrow and a signed number. Up means more risk, so up is worse.
- **Also changed**: the overall score, the count of Critical and High domains, decisions taken in this session, average time to fix, and missed deadlines.
- Arrows and signs as well as colours, so the page still reads in black and white print or for colour-blind readers.

Domain	June	Now	Change
AI Security and Governance	51.0	53.5	▲ +2.5 (worse)
Ransomware Readiness	57.6	55.2	▼ −2.4 (better)
Attack Surface and VAPT	74.0	76.0	▲ +2.0 (worse)
SOC / MDR	70.0	72.0	▲ +2.0 (worse)
Third-Party Risk	63.0	65.0	▲ +2.0 (worse)

Overall score moved from 51.0 to 50.8. The count of Critical domains stayed at 2 and High at 2. Average time to fix fell from 41 to 36 days, and missed deadlines fell from 9 to 7 (both illustrative).

Real-time example: the board meeting. A board member asks, "Are we better than June?" Sara opens this tab and answers in three sentences: "The headline is the same, High. Ransomware readiness improved by 2.4 points after the recovery tests, but attack surface, SOC and third-party risk each got about 2 points worse. We are fixing things faster, 36 days against 41, and missing fewer deadlines." The conversation shifts from "is the number right" to "why are those three still drifting up".

Watch out for. Moves of 2 points are small. Read this tab together with the 6-month trend lines on the Domain posture tab to tell a blip from a steady drift, as with Attack Surface, which has crept up from 70 to 76 over six months.

8. Board pack: "What do I send the board?"

This tab writes a draft board paper from the live data, and every sentence in it names the data table it came from.

In plain words. Preparing the quarterly cyber paper usually means days of copying figures from screens into slides, and nobody can later say where a number came from. Here the paper is assembled from a fixed template. Each sentence carries a small source tag in brackets, such as [enterprise_scores_v2] or [loss_models], so a reader can trace it back. A named person must **review and approve** it before it goes out. DWAN drafts; people decide.

What you see on the screen

- Buttons: **Export to PDF / print** and **Mark reviewed**.
- A page headed "Board pack, Al Dhabi Holdings, snapshot 5 Sep 2026, DRAFT (sample data)".
- Five blocks: the headline, the money, what changed since June, decisions needed, and appetite and accepted risks, followed by a confidence statement.

What the draft says, in the sample

- **Headline:** "Cyber risk is High, driven by 2 Critical domains". It explains that the score is 50.8 (Medium) and that the floor rule raises it to High, with the spread of 2 Critical, 2 High, 4 Medium and 4 Low.
- **Money:** AED 2.9M expected a year and about AED 10.1M in a bad year (1 in 10). The largest driver is ransomware outage at AED 1.78M expected.

- **What changed:** the biggest movers since June.
- **Decisions needed:** the four items from the Decision queue, each with its recommendation.
- **Appetite and accepted risks:** 5 of 12 domains outside appetite, and the HR portal acceptance that needs renewing.
- **Confidence:** medium (72%), with the three weakest domains named: AI Security (25%), OT / ICS (40%) and Third-Party Risk (60%).

Real-time example: Thursday before the board. The board meets on Friday morning. On Thursday afternoon Sara opens this tab. The draft already exists, and it matches the numbers she saw on tabs 1 to 7. She checks each statement against its source tag, softens one sentence, and presses **Mark reviewed**. Then she exports a PDF. What usually means days of copying figures into slides becomes one review pass, and when a director asks "where does the 10.1 million come from?", the answer is one tag away.

Watch out for. The pack is only as good as the data behind it, so the confidence statement is deliberately part of it. A board that reads "confidence 72%" understands the numbers are good, not perfect.

One week with the tabs, and a glossary

Sara's week uses each tab for a different job, and most of the work happens on Monday and Thursday.

When	Tab	What Sara does
Monday, 8:30	1. Board headline	Reads the verdict and the money, briefs the CEO in two minutes
Monday, 9:00	4. Decision queue	Approves the payment gateway patch window, sets the 7-day clock
Tuesday	3. Trust panel	Asks the team to restart the stale connectors and close coverage gaps
Wednesday	5. What-if	Tests recovery testing and immutable backup before the budget meeting
Thursday	6. Closure and acceptance	Runs verification scans, renews or revokes the expiring HR portal permit

Thursday	2. Domain posture and 7. What changed	Checks which domains are drifting and why
Thursday	8. Board pack	Reviews the draft, marks it reviewed, exports the PDF

Glossary in plain words

- **Domain**: one area of security, such as email, cloud or ransomware recovery. There are 12.
- **Score (0 to 100)**: how much risk is left in a domain. Higher is worse.
- **Band**: Low (under 35), Medium (35 to 54), High (55 to 69), Critical (70 and above).
- **Coverage**: the share of your assets that the security tools actually watch.
- **Adjusted score**: the score after assuming that anything unwatched is at least as risky as a default of 60.
- **Confidence**: how much to trust a number, from fresh feeds and coverage.
- **Feed**: one tool's regular delivery of data into DWAN. It is "stale" when it is overdue.
- **Floor rule**: any Critical domain stops the headline reading better than High.
- **Appetite**: the most risk the board says it will tolerate in a domain.
- **Expected loss (ALE)**: the modelled average cost per year of cyber events. A **bad year** figure (1 in 10, or 1 in 20) shows a rare large event.
- **Crown jewel**: an asset the business cannot afford to lose, such as the payment gateway.
- **KEV and EPSS**: a list of weaknesses already being attacked, and a 0 to 1 estimate of how likely an attack is soon.
- **SLA**: the deadline for fixing a finding.
- **Verified closure**: a fix that a fresh scan has confirmed.
- **Risk acceptance**: a knowing decision to live with a risk, with a reason, an approver and an expiry date.
- **Floor, appetite and weights** are set in the scoring policy, which the CISO approves. Approved versions are frozen and changes are logged.

All names, figures and dates in this guide come from the prototype's sample data for Al Dhabi Holdings. Sara and her Monday are illustrations, not a real client.